

WATERLINE

PROOF OF DELIVERED COMPUTE

Check that the GPU you rent delivers what you pay for: the right chip, at its rated speed.

THE PROBLEM

Compute is becoming a financial asset. Its delivery is still self-reported.

~\$400B

GPU cloud by 2031, up from \$25B in 2025

Synergy Research

\$8.5B

lent against GPUs in one facility, rated investment grade

CoreWeave

Oct 5

H100 rental futures scheduled to trade on CME

CME Group

1,755 → 345 MHz

an H100 under load, billed at full price
`matt.sh`

There is no independent, shared record.

01

Sealed work, a deadline, a core count.

01

A fresh problem

Seeded INT8 matrices, 8.8 trillion operations a step, on the tensor cores. Nothing to precompute.

02

A clock on our side

The deadline runs on the API, sized to the listed GPU's rated speed.

03

Sealed, then re-graded

The GPU commits every answer first. Then we pick 8 rows and recompute them on a CPU.

04

Count the cores

A timing staircase: 132 on an H100 SXM, 108 on an A100. Heat can't remove cores.

PASS

The listed chip, on time. Published at once.

DEGRADED

The listed chip, too slow. Published with its numbers.

FAIL

Another chip, or wrong answers. Published with the listing.

EVERY FAILURE CARRIES ITS LISTING

A pass needs real silicon. A failure needs the listing you rented.

The listing

- The reporter pastes the listing they rented
- Read automatically: a different GPU stops the report
- Shown beside the failure for anyone to compare

Two reports to fail

- One report: suspect. Two: failed, on the GPU and its provider
- Suspicious reports are flagged automatically
- Two passes after a failure: recovered

Each report is its own vote

Every GPU has a name. Every provider keeps a score.

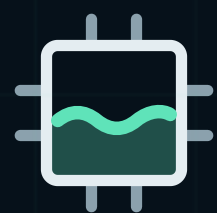


waterline.eth

```
|─ cloud-a.waterline.eth 0 of 2 failed
|   └─ gpu-... pass · H100 SXM
└─ cloud-b.waterline.eth 1 report
    └─ gpu-... suspect · 1 of 2
```



- Builds every write; we sign it
- Indexes every verdict, calls our webhook
- The agent skips failed GPUs and providers: the record decides



Proof of Delivered Compute.

Any renter can run the check. The result stays on the GPU's name for whoever needs it next: a renter, a lender, an agent, or even a GPU credit provider.



ENSv2



Curvegrid MultiBaas

waterline-eth.vercel.app

github.com/FlyingWHR/waterline